

Name

CWID

# Quiz 1

CS525

Advanced Database Organization  
Spring 2021

Due: March 25<sup>th</sup> 11:59pm

Results

---

Please leave this empty! 1.1.1  1.2.1  1.2.2  1.2.3  Sum

# Instructions

- You have to hand in the assignment using your blackboard
- This is an individual and not a group assignment. Fraud will result in 0 points
- For your convenience the number of points for each part and questions are shown in parenthesis.
- There are **2** parts in this quiz
  1. Disk Organization
  2. Index Structures

## Part 1.1 Disk Organization (Total: 18 Points)

### Question 1.1.1 Disk Access (18 Points)

Consider a disk with a sector size of 512 bytes, 2000 tracks per surface, 50 sectors per track, five double-sided platters, and average seek time of 10 msec. Suppose that a block size of 1024 bytes is chosen. Suppose that a file containing 100,000 records of 100 bytes each is to be stored on such a disk and that no record is allowed to span two blocks.

[3 Points each]

- What is the capacity of a track in bytes? What is the capacity of each surface? What is the capacity of the disk?

#### Solution

```
bytes/track = bytes/sector × sectors/track = 512 × 50 = 25,600 bytes ≈ 25K  
bytes/surface = bytes/track × tracks/surface = 25,600 × 2000 = 51,200,000 bytes ≈ 50,000K  
bytes/disk = bytes/surface × surfaces/disk = 51,200,000 × 5 × 2 = 512,000,000 bytes ≈ 500,000K
```

- How many records fit onto a block?

#### Solution

$\left\lfloor \frac{\text{Block size}}{\text{Record size}} \right\rfloor = \left\lfloor \frac{1024}{100} \right\rfloor = 10$ . We can have at most 10 records in a block.

- How many blocks are required to store the entire file?

#### Solution

```
File size = # Record × Record size  
File size = 100,000 × 100 = 10,000,000 bytes  
We need  $\lceil \frac{\text{File size}}{\text{Block size}} \rceil = \lceil \frac{10,000,000}{1024} \rceil = 9765.625 \approx 10,000$  blocks to store the file.
```

- If the file is arranged sequentially on the disk, how many surfaces are needed?

#### Solution

```
# Sectors/Tracks = 50, Sector size = 512 bytes, Block size = 1024 bytes  
One track has 25 blocks,  
One cylinder has  $25 \times 5 \times 2 = 250$  blocks.  
We need 10,000 blocks to store this file.  
So, we need  $\lceil \frac{10,000}{250} \rceil = 40$  cylinders  
We need 10 surfaces to store the file.
```

- How many records of 100 bytes each can be stored using this disk?

#### Solution

The capacity of the disk is 512,000,000 bytes ≈ 500,000K, which has ≈ 500,000 blocks.  
Each block has 10 records.  
The disk can store no more than 512,000,000 ≈ 5,000,000 records.

- If pages are stored sequentially on disk, with page 1 on block 1 of track 1, what page is stored on block 1 of track 1 on the next disk surface?

#### Solution

There are 25 blocks in each track.  
It is block 26 on block 1 of track 1 on the next disk surface.

## Part 1.2 Index Structures (Total: 54 Points)

### Question 1.2.1 B<sup>+</sup>-tree Construction (10 Points)

Assume that you have the following table:

Item		
SSN	name	age
2	Pete	13
3	Bob	23
29	Heinz	14
5	John	49
19	Manny	33
23	Gertrud	29
11	Alice	77
17	Lily	3
39	Sammy	34
7	Joe	45

Create a B<sup>+</sup>-tree for table **Item** on key *SSN*. Assume that the tree is initially empty and values are added in the order shown in the table above. Construct B<sup>+</sup>-tree for the cases where the **number of pointers that will fit in one node** is as follows:

- a. Three
- b. Six

Write down the resulting B<sup>+</sup>-tree after each step and when splitting or merging nodes follow these conventions:


- **Leaf Split:** In case a leaf node needs to be split during insertion and  $n$  is even, the left node should get the extra key. E.g, if  $n = 2$  and we insert a key 4 into a node [1,5], then the resulting nodes should be [1,4] and [5]. For odd values of  $n$  we can always evenly split the keys between the two nodes. In both cases the value inserted into the parent is the smallest value of the right node.
- **Non-Leaf Split:** In case a non-leaf node needs to be split and  $n$  is odd, we cannot split the node evenly (one of the new nodes will have one more key). In this case the “middle” value inserted into the parent should be taken from the right node. E.g., if  $n = 3$  and we have to split a non-leaf node [1,3,4,5], the resulting nodes would be [1,3] and [5]. The value inserted into the parent would be 4.
- **Node Underflow:** In case of a node underflow you should first try to redistribute values from a sibling and only if this fails merge the node with one of its siblings. Both approaches should prefer the left sibling. E.g., if we can borrow values from both the left and right sibling, you should borrow from the left one.

## Solution


a. Three  
Insert 2, 3

• 2 • 3 •


Insert 29, 5




Insert 5, 19




Insert 23




Insert 11



Insert 17, 39




Insert 7




b. Six


Insert 2, 3, 29, 5, 19



Insert 23, 11, 17, 39



Insert 7




### Question 1.2.2 Operations (20 Points)

Given is the B<sup>+</sup>-tree shown below ( $n = 3$ ). Execute the following operations and write down the resulting B<sup>+</sup>-tree after each operation:


**insert(3), delete(5), insert(10), insert(8), delete(29), delete(19), delete(11)**

Use the conventions for splitting and merging introduced in the previous question.




**Solution**


insert(3), delete (5), insert (10)




insert(8)



delete(29), then delete(19)



delete(11)



### Question 1.2.3 Extendable Hashing-tree Construction (24 Points)

Suppose that we are using extendable hashing on a file that contains records with the following search key values: 7, 15, 44, 11, 19, 3, 33, 17, 14, 43, 18, 20

Show the extendable hash structure for this file if the hash function is  $h(x) = x \bmod 7$  and buckets can hold 2 records

#### Solution


Initial Index Table



Insert 7, 15




Insert 44, 11, 19, 3




## Solution

Insert 33




Insert 17



## Solution

Insert 14



## Solution

Insert 43, 18, 20

